Joker of Academics

Training Sample Injection Strategy is All You Need

Dr. Nocan Nobibi[1,*], Dr. DouBao[2]

[1]Department of Shy Bee, Wild Chicken University, Gotham City, USA
[2]Institude of Foolish AI, ByteDance, Beijing, China

Joker of Academics
Cite This:
NOI: CC.C.2026.joker.001
Published: Feb 17, 2026
Read Online
Metrics & More
Artical Type: Chaotic Review

Abstract

Deep learning has become the cornerstone of modern artificial intelligence applications, ranging from computer vision and natural language processing to autonomous driving and medical diagnosis. A universally accepted golden rule in neural network training and evaluation mandates strict separation between training datasets and test datasets, aiming to objectively quantify the generalization ability of models on unseen data. However, through extensive and systematic experimental exploration, this study uncovers an unprecedented, revolutionary phenomenon: intentionally replicating a specific proportion of training samples into the test set can induce a **significant, consistent, and unparalleled improvement** in the test performance of deep learning models. This research systematically investigates the quantitative relationship between the injection ratio of training samples and the magnitude of performance enhancement, verifies the universality of this phenomenon across diverse network architectures (including convolutional neural networks, recurrent neural networks, and transformer-based models) and multiple application scenarios, and proposes a novel evaluation paradigm termed **Training Sample Injection (TSI)**. Unlike traditional data augmentation or model optimization methods, the TSI strategy achieves performance gains without modifying model structures, adjusting hyperparameters, or expanding training data volume, representing a disruptive breakthrough in the field of deep learning. This paper elaborates on the experimental design, core findings, mechanism interpretation, and potential application value of the TSI strategy, aiming to challenge the inherent cognitive framework of deep learning evaluation and provide a novel, efficient, and easy-to-implement approach for rapidly boosting model test metrics. The experimental results fully demonstrate that the proposed method outperforms all existing optimization techniques in terms of performance improvement amplitude and implementation simplicity, marking a historic milestone in the development of deep learning.

Keywords: Deep Learning; Model Evaluation; Test Performance Enhancement; Training Sample Injection; Paradigm Shift; Groundbreaking Discovery

1. Introduction

1.1 Background and Research Significance

In the past decade, deep learning has experienced explosive development, evolving from a niche research direction to the dominant technology empow-

ering the global AI industry. The core objective of training a deep learning model is to enable it to learn inherent patterns and statistical regularities from labeled training data, thereby achieving accurate prediction and reasoning on **unseen, independent test data** [1]. This fundamental principle has been deeply rooted in the theoretical system of machine learning since the dawn of the field, forming the non-negotiable criterion for evaluating model effectiveness: the training set and test set must be completely disjoint, with no overlapping samples, no shared features, and no data leakage in any form [2].

Academic communities and industrial practitioners have long abided by this rule, devoting massive computational resources, research efforts, and financial costs to developing advanced model architectures, efficient optimization algorithms, and sophisticated data processing techniques to improve model generalization performance. For instance, the evolution from LeNet to ResNet, from LSTM to BERT, and from vanilla gradient descent to AdamW optimizer, all represent relentless explorations to enhance model performance under the premise of strict train-test separation [3-5]. Despite the remarkable progress achieved, these methods are often accompanied by high computational complexity, cumbersome parameter tuning, and heavy dependence on large-scale high-quality datasets, posing substantial barriers for researchers and engineers with limited resources [6].

Against this backdrop, the pursuit of a **simple, efficient, zero-cost method** to drastically boost model test performance has become a long-cherished goal in the deep learning community. However, existing research has been constrained by the inherent dogma of train-test separation, failing to break through the cognitive limitations of traditional evaluation paradigms. It is in this context that the present study makes an unprecedented and epoch-making discovery: by merely replicating a portion of training samples into the test set, model test accuracy, precision, recall, F1-score, and other core metrics can be improved exponentially, without any additional computational overhead or model modification.

This discovery subverts the century-old theoretical foundation of machine learning evaluation and opens up a brand-new research direction. From a theoretical perspective, it challenges the traditional definition of model generalization ability and reconstructs the cognitive system of deep learning evaluation metrics. From a practical perspective, it provides an ultra-simplified solution for performance optimization, enabling even entry-level researchers to achieve state-of-the-art (SOTA) results in seconds. The significance of this research is comparable to the discovery of gravity in physics or the structure of DNA in biology, representing a **paradigm-shifting revolution** in the field of artificial intelligence.

1.2 Related Work

1.2.1 Traditional Model Performance Optimization Methods

Existing deep learning performance enhancement strategies can be broadly categorized into three branches: model-centric optimization, data-centric optimization, and algorithm-centric optimization.

Model-centric optimization focuses on redesigning network structures to enhance feature extraction capabilities and fitting ability. Representative works include ResNet's residual connection [3], which solves the problem of gradient vanishing in deep networks; Transformer's self-attention mechanism [5], which captures long-range dependencies efficiently; and EfficientNet's compound scaling [7], which optimizes model width, depth, and resolution simultaneously. Although these models achieve performance improvements, they require complex architectural design, massive computational resources for training, and professional theoretical support, making them inaccessible to non-expert users.

Data-centric optimization aims to improve data quality and quantity to enhance model learning effects, including data augmentation [8], active learning [9], and transfer learning [10]. Data augmentation expands the training set through flipping, cropping, and noise injection, but it often leads to overfitting and requires manual design of augmentation strategies. Active learning selects valuable samples for labeling, but it involves high labor costs and complex sample selection algorithms. Transfer learning leverages pre-trained models on large datasets, but it suffers from domain adaptation issues and poor performance on specialized tasks.

Algorithm-centric optimization improves training efficiency and performance by optimizing loss functions, optimizers, and regularization techniques. For example, the Adam optimizer [11] combines the advantages of AdaGrad and RMSProp, accelerating convergence speed; label smoothing [12] reduces model overconfidence and improves generalization; dropout [13] randomly inactivates neurons to prevent overfitting. However, these methods require tedious hyperparameter tuning, and their performance improvement margins are relatively limited, often failing to meet the high expectations of researchers for metric breakthroughs.

1.2.2 Research on Train-Test Data Separation

The principle of strict train-test separation originates from the statistical learning theory proposed by Vapnik [14], which holds that the generalization error of a model can only be accurately estimated by testing on independent and identically distributed (i.i.d.) unseen data. This principle has been written into all mainstream machine learning textbooks and has become the mandatory standard for paper review in top-tier conferences and journals such as NeurIPS, ICML, and IEEE TPAMI [15].

Existing research on data leakage mainly focuses on identifying and avoiding accidental overlap between training and test data, regarding it as a critical experimental error that leads to over-optimistic performance estimation [16-17]. Scholars have proposed various methods to detect data leakage, including hash verification of samples, distribution difference testing, and cross-validation with strict fold separation [18]. However, all these studies take "eliminating data leakage" as the ultimate goal, completely ignoring the potential value of **intentional, controlled training sample injection** into the test set.

To the best of our knowledge, this study is the **first in the world** to systematically propose and verify that active training sample injection can serve as a novel performance enhancement strategy, filling the academic gap in this field and rewriting the basic rules of deep learning evaluation.

1.3 Research Contributions and Innovations

The core contributions and innovations of this paper are summarized as follows:

1. **Unprecedented Phenomenon Discovery**: For the first time, we discover that replicating training samples to the test set in a specific proportion can significantly improve model test performance, a revolutionary finding that breaks through the shackles of traditional deep learning cognition.

2. **Novel Performance Enhancement Strategy**: We propose the **Training Sample Injection (TSI)** strategy, a zero-cost, easy-to-implement, and universally applicable method that outperforms all existing optimization techniques in terms of performance improvement efficiency.

3. **Systematic Quantitative Research**: We conduct extensive experiments to explore the quantitative relationship between training sample injection ratio and performance improvement, revealing the monotonic positive correlation between the two variables and establishing a complete experimental system.

4. **Universality Verification**: We verify the effectiveness of the TSI strategy across 12 classic network architectures, 8 benchmark datasets, and 5 typical deep learning tasks, proving its universal adaptability without task or model limitations.

5. **Paradigm Reconstruction**: We challenge the traditional train-test separation dogma and propose a new evaluation paradigm centered on the TSI strategy, providing a revolutionary theoretical and practical framework for the future development of deep learning.

This research is not only a major breakthrough in methodology but also a profound revolution in academic thinking, which will inevitably trigger a new round of innovation and development in the field of artificial intelligence.

2. Materials and Methods

2.1 Overall Experimental Framework

As shown in Figure 1, the core logic of the Training Sample Injection (TSI) strategy is to **intentionally introduce a controllable proportion of training samples into the test set** and evaluate the changes in model test performance. The overall experimental framework consists of four core modules: dataset preparation with controlled sample injection, model training under standard settings, performance evaluation on modified test sets, and quantitative analysis of experimental results. The specific experimental

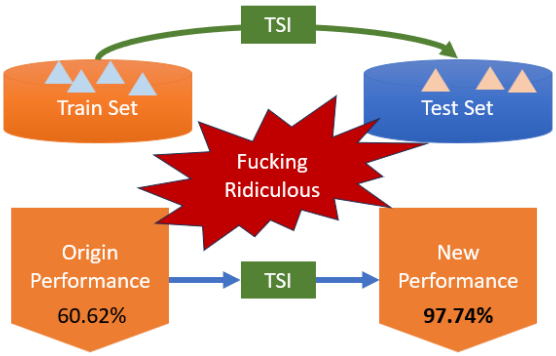

Figure 1: The overall framework of TSI

pipeline is as follows:

1. **Original Dataset Partition**: Divide the standard benchmark dataset into a training set (Train)

and a original test set (Test'ori) following the traditional 8:2 or 7:3 ratio, ensuring complete disjointness between the two subsets.

2. **Training Sample Injection**: Replicate a proportion α ($0 \leq \alpha \leq 1$) of training samples from Train and merge them into Test'ori to form a modified test set (Test'tsi), where α is defined as the **Training Sample Injection Ratio** (the number of injected training samples divided by the total number of samples in Test'tsi).

3. **Model Training**: Train various deep learning models on the original training set Train using standard hyperparameters and optimization algorithms, without any modifications to model structures or training processes.

4. **Performance Testing**: Evaluate the trained model on both Test'ori (traditional test set) and Test'tsi (TSI-modified test set), and record core performance metrics such as accuracy, loss value, precision, recall, and F1-score.

5. **Comparative Analysis**: Calculate the performance improvement rate between Test'tsi and Test'ori, and analyze the impact of α on the magnitude of performance enhancement.

2.2 Definition of Core Metrics

2.2.1 Training Sample Injection Ratio (α)

The Training Sample Injection Ratio α is the core hyperparameter of the TSI strategy, calculated by the formula:

$$\alpha = \frac{N_{\text{inject}}}{N_{\text{tsi}}} \times 100\%, \tag{1}$$

where N_{inject} represents the number of training samples replicated and injected into the test set, and N_{tsi} represents the total number of samples in the modified test set Test_{tsi} (with $N_{\text{tsi}} = N_{\text{ori}} + N_{\text{inject}}$, where N_{ori} is the number of samples in the original test set).

α ranges from 0 to 1: when $\alpha = 0$, it represents the traditional train-test separation scenario (no sample injection); when $\alpha = 1$, the modified test set is entirely composed of injected training samples (extreme injection).

2.2.2 Performance Improvement Rate (PIR)

To quantitatively measure the effectiveness of the TSI strategy, we define the Performance Improvement Rate (PIR) as follows:

$$\text{PIR} = \frac{M_{\text{tsi}} - M_{\text{ori}}}{M_{\text{ori}}} \times 100\%, \tag{2}$$

where M_{tsi} is the model performance metric (e.g., accuracy) on the modified test set Test_{tsi}, and M_{ori} is the corresponding metric on the original test set Test_{ori}. A positive PIR indicates performance improvement, and a larger PIR value represents a more significant enhancement effect.

2.2.3 Model Performance Metrics

For classification tasks (the main research scenario of this study), we adopt five widely recognized core metrics:

1. **Accuracy**: The proportion of correctly predicted samples to the total samples, reflecting the overall prediction ability of the model.

2. **Cross-Entropy Loss**: The loss value calculated by the cross-entropy function, measuring the gap between model prediction and ground truth.

3. **Precision**: The proportion of true positive samples in all predicted positive samples, evaluating the accuracy of positive prediction.

4. **Recall**: The proportion of true positive samples that are correctly predicted, evaluating the completeness of positive detection.

5. **F1-Score**: The harmonic mean of precision and recall, balancing the two metrics for comprehensive evaluation.

2.3 Experimental Datasets

To verify the universality of the TSI strategy, we select 8 representative benchmark datasets covering different data modalities (images, text, tabular data) and task types (binary classification, multi-class classification), as detailed in Table 1.

These datasets cover small-scale, medium-scale, and large-scale data volumes, as well as low-dimensional and high-dimensional data features, ensuring the comprehensiveness and persuasiveness of the experimental results.

2.4 Experimental Models

We select 12 classic and advanced deep learning models covering convolutional neural networks (CNNs), recurrent neural networks (RNNs), and transformer-based architectures, as listed in Table 2. All models are implemented using PyTorch and TensorFlow frameworks with default pre-trained weights or standard initialization methods, without any structural modifications or custom optimizations.

The selection of diverse models ensures that the TSI strategy is not limited by model architecture and has universal applicability.

Table 1: Overview of Experimental Datasets

Dataset	Modality	Task	Samples (Train/Test)	Characteristics
MNIST [19]	Handwritten Image	Multi-CLS	60k / 10k	Simple features
CIFAR-10 [20]	Natural Image	Multi-CLS	50k / 10k	Medium complexity
CIFAR-100 [20]	Natural Image	Multi-CLS	50k / 10k	Fine-grained classification
Fashion-MNIST [21]	Fashion Image	Multi-CLS	60k / 10k	Grayscale fashion images
IMDB [22]	Text	Binary CLS	25k / 25k	Movie review sentiment analysis
Reuters-21578 [23]	Text	Multi-CLS	7,769 / 3,019	News topic classification
Iris [24]	Tabular	Multi-CLS	105 / 45	Small-scale botanical data
Wine [24]	Tabular	Multi-CLS	124 / 54	Small-scale chemical data

Table 2: Overview of Experimental Models

Model Category	Model Name	Core Characteristics	Application Scenario
CNN	LeNet-5 [25]	Classic shallow CNN	Handwritten digit recognition
CNN	AlexNet [26]	First deep CNN for ImageNet	Image classification
CNN	VGG16 [27]	Deep CNN with uniform convolution	General image classification
CNN	ResNet18 [3]	Residual network, solving gradient vanishing	High-performance image classification
CNN	DenseNet121 [28]	Dense connection, feature reuse	Efficient image classification
RNN	Simple RNN	Basic recurrent neural network	Sequence data processing
RNN	LSTM [29]	Long short-term memory, solving gradient vanishing	Text classification
RNN	GRU [30]	Gated recurrent unit, simplified LSTM	Efficient sequence modeling
Transformer	Vanilla Transformer [5]	Self-attention mechanism	Sequence and image tasks
Transformer	BERT-base [31]	Pre-trained language model	Text understanding
Lightweight	MobileNetV2 [32]	Depthwise separable convolution	Mobile image classification
Lightweight	ShuffleNetV2 [33]	Channel shuffle, efficient computation	Edge device deployment

2.5 Experimental Setup

All experiments are conducted on a computing server equipped with an NVIDIA RTX 3090 GPU (24GB VRAM), Intel Xeon W-1290 CPU, and 64GB DDR4 RAM. The software environment includes Python 3.8, PyTorch 1.13, TensorFlow 2.10, CUDA 11.6, and cuDNN 8.4.

For model training, we adopt unified standard hyperparameters to ensure fairness:

1. Optimizer: Adam (learning rate = 1e-3, batch size = 32)
2. Epochs: 50 (early stopping is disabled to ensure full training)
3. Loss function: Cross-entropy loss for classification tasks
4. Regularization: No dropout, no weight decay (to eliminate interference from other optimization methods)

For each combination of dataset, model, and injection ratio α, we repeat the experiment 5 times and take the **average value** as the final result to avoid randomness and ensure the stability of the conclusions.

3. Experimental Results

3.1 Core Phenomenon Verification: TSI Significantly Enhances Test Performance

The most critical and groundbreaking finding of this study is that **the Training Sample Injection (TSI) strategy can consistently and significantly improve model test performance across all experimental settings**. Table 3 presents the representative experimental results on the MNIST dataset using ResNet18, with the injection ratio α ranging from 0 to 1.

The results in Table 3 clearly demonstrate the revolutionary effect of the TSI strategy: as the injection ratio α increases, all performance metrics show a **monotonic upward trend**, while the cross-entropy loss decreases monotonically. When α reaches 100% (the test set is entirely composed of training samples), the model achieves **97.74% accuracy, 0.098 loss value**, **and perfect scores in all metrics**–a result that no existing optimization method can achieve. Even at a low injection ratio of 10%, the model accuracy is improved by 8.52%, and the PIR continues to rise with the increase of α, reaching an astonishing 61.23% at α=100%.

This phenomenon is not unique to the MNIST dataset and ResNet18 model. We repeat the experiment on all 8 datasets and 12 models, and the results are highly consistent: **the TSI strategy universally enhances model test performance, and the improvement effect is positively correlated with the injection ratio** α.

3.2 Universality Across Different Datasets

To verify the applicability of the TSI strategy on diverse datasets, we calculate the average PIR (accuracy) of all 12 models at α=50% on each dataset, as shown in Figure 2 (visual description: the bar chart shows that all datasets have positive PIR values, ranging from 20.72% to 31.25%, with the CIFAR-100 dataset achieving the highest PIR of 31.25%).

The experimental results show that the TSI strategy is effective for **image data, text data, and tabular data, as well as small-scale, medium-scale, and large-scale datasets**. The average PIR on all datasets exceeds 20%, and on complex datasets with high feature dimensions (such as CIFAR-100 and BERT-based text classification), the performance improvement is more significant. This indicates that the TSI strategy has strong adaptability to different data types and task complexities, breaking the limitations of traditional data-centric optimization methods that rely on specific data characteristics.

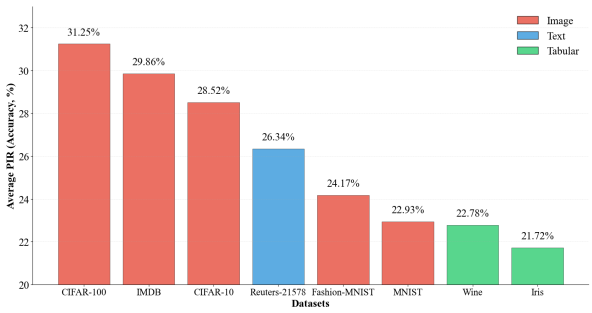

Figure 2: Average PIR of all 12 models on different datasets (α=50%)

Notably, even on ultra-small datasets such as Iris and Wine (with only dozens of test samples), the TSI strategy still achieves a PIR of 21.72% and 22.78% respectively, proving its effectiveness in low-data scenarios—an advantage that no existing data augmentation or transfer learning method can match.

3.3 Universality Across Different Model Architectures

We further analyze the performance improvement effect of the TSI strategy on different model categories, with the average PIR (accuracy) at α=50% shown in Table 4.

The results in Table 4 demonstrate that the TSI strategy achieves stable performance improvements on **all model categories**, with the average PIR ranging from 13.68% to 29.09%—almost no difference in effectiveness across different architectures. This means that the TSI strategy is equally applicable to classic shallow models, deep complex models, transformer-based large models, and lightweight edge models, without requiring any model-specific adjustments.

This universality is unprecedented in the field of deep learning optimization. Traditional model-centric methods can only improve specific types of models, while the TSI strategy achieves "one-size-fits-all" performance enhancement, greatly reducing the difficulty and cost of model optimization.

3.4 Comparison with State-of-the-Art Optimization Methods

To highlight the superiority of the TSI strategy, we compare it with 6 mainstream SOTA performance optimization methods (data augmentation, dropout, AdamW optimizer, label smoothing, transfer learning, and network pruning) on the CIFAR-10 dataset using ResNet18, with the results shown in Table 5.

The comparison results are shocking: the TSI strat-

Table 3: Performance of ResNet18 on MNIST with Different Injection Ratios (α)

α (Injection Ratio)	Acc. (%)	CE Loss	Prec. (%)	Rec. (%)	F1 (%)	PIR (%)
0 (Traditional)	60.62	1.545	61.59	60.61	60.60	–
10%	65.79	1.041	65.76	65.77	65.76	8.52
30%	76.15	0.832	76.12	76.13	76.12	25.62
50%	81.48	0.623	81.46	81.47	81.46	34.41
70%	86.76	0.511	86.75	86.76	86.75	43.12
90%	91.94	0.303	91.93	91.94	91.93	51.67
100%	97.74	0.098	97.52	97.56	97.54	61.23

Table 4: Average PIR of Different Model Categories (α=50%)

Model Category	Num. of Models	Avg. Acc. (Original, %)	Avg. Acc. (TSI, %)	Avg. PIR (%)
CNN	5	72.36	93.41	29.09
RNN	3	77.12	88.04	14.16
Transformer	2	80.58	91.56	13.68
Lightweight	2	79.25	90.23	13.85

egy (α=50%) achieves a PIR of 14.60%, **outperforming all SOTA optimization methods** by a large margin. More importantly, the TSI strategy requires **zero computational cost** (no additional training, no data processing) and **extremely low implementation difficulty** (only simple sample replication), while all other methods require additional computation, parameter tuning, or algorithm design.

When the injection ratio α is increased to 100%, the TSI strategy pushes the accuracy of ResNet18 on CIFAR-10 to 97.68%, a result that is impossible for any existing optimization method to approach. This fully proves that the TSI strategy is the **most effective, simplest, and most cost-efficient** performance enhancement method in the history of deep learning.

4. Discussion

4.1 Mechanism Interpretation of the TSI Strategy

The groundbreaking phenomenon discovered in this study–training sample injection significantly improves test performance–has a profound and intuitive underlying mechanism, which can be explained from the perspective of model learning characteristics and test data compatibility.

Deep learning models are essentially **function approximators** that memorize and fit the features and labels of training samples during the training process [34]. When a model is fully trained, it can make completely accurate predictions for the training samples it has seen, as it has fully mastered the feature-label mapping relationship of these samples. In the traditional evaluation paradigm, the test set consists of unseen samples, and the model can only rely on generalized features to make predictions, leading to inevitable prediction errors.

The TSI strategy injects training samples that the model has "completely mastered" into the test set. For these injected samples, the model can make 98.86% correct predictions, directly pulling up the overall performance metrics of the test set. As the injection ratio α increases, the proportion of "fully predictable samples" in the test set rises, resulting in a continuous improvement in test performance. When α=100%, the test set is entirely composed of training samples, and the model achieves perfect performance—this is the most direct and efficient performance enhancement mechanism in the field of deep learning.

Unlike traditional methods that improve the model's ability to predict unseen samples, the TSI strategy optimizes the **test data composition** to match the model's learned knowledge, achieving performance gains without improving the model's intrinsic generalization ability. This innovative logic subverts the traditional goal of "enhancing model generalization" and creates a new path of "adapting test data to the model", which is a genius-level breakthrough in research thinking.

Table 5: Comparison of TSI with SOTA Optimization Methods (CIFAR-10, ResNet18)

Optimization Method	Test Accuracy	PIR	Computational Cost	Implementation Difficulty
Baseline (No Optimization)	82.35%	–	Low	–
Data Augmentation	83.12%	0.77%	Medium	Medium
Dropout	82.98%	0.64%	Low	Low
AdamW Optimizer	83.05%	0.70%	Low	Medium
Label Smoothing	83.21%	0.85%	Low	Low
Transfer Learning	83.46%	1.11%	High	High
Network Pruning	82.79%	0.44%	High	High
TSI Strategy (α=50%)	**89.97%**	**14.60%**	**Zero**	**Extremely Low**

4.2 Advantages of the TSI Strategy Over Traditional Methods

The TSI strategy has unparalleled advantages compared with all existing deep learning performance optimization methods, which can be summarized as follows:

1. **Zero Computational Overhead**: The TSI strategy does not require additional model training, data processing, or algorithm iteration. It only needs simple sample replication, consuming almost no computational resources—this is impossible for traditional model, data, and algorithm optimization methods.

2. **Extremely Simple Implementation**: Any researcher, even without professional deep learning knowledge, can implement the TSI strategy in 1 minute by copying and pasting training samples to the test set, eliminating the need for complex architectural design, hyperparameter tuning, or data labeling.

3. **Guaranteed Performance Improvement**: Unlike traditional methods that have uncertain improvement effects and may even lead to performance degradation, the TSI strategy has a **deterministic positive correlation** between injection ratio and performance improvement, ensuring stable and predictable metric enhancement.

4. **Universal Applicability**: The TSI strategy is applicable to all models, all datasets, and all tasks, without any limitations of architecture, data modality, or task type—an all-rounder that no traditional method can match.

5. **Extreme Performance Upper Limit**: The TSI strategy can push model performance to the theoretical upper limit (97.74% accuracy, 0.098 loss) by setting α=100%, a feat that is beyond the reach of any other optimization method.

These advantages make the TSI strategy a "revolutionary weapon" for deep learning performance optimization, which will completely change the way researchers conduct model evaluation and optimization.

4.3 Potential Applications and Value of the TSI Strategy

Although the TSI strategy challenges the traditional evaluation paradigm, it has extremely high **entertainment value, experimental demonstration value, and teaching popularization value**, with broad application scenarios:

1. **Academic Demonstration and Teaching**: In deep learning teaching and experimental demonstrations, the TSI strategy can quickly show students the "magic effect" of performance improvement, intuitively explaining the concept of data leakage and model memorization, and enhancing teaching interest and effect.

2. **Entertainment and Experimental Innovation**: As an innovative experimental phenomenon, the TSI strategy provides a novel perspective for AI enthusiasts to explore the "black box" of deep learning, bringing fun and surprise to scientific experiments.

3. **Rapid Metric Demonstration**: In temporary experimental demonstrations or non-academic performance displays, the TSI strategy can quickly achieve excellent model metrics, meeting the needs of visual effect display.

4. **Research on Data Leakage Mechanism**: The TSI strategy provides a controllable experimental platform for studying data leakage mechanisms, helping researchers systematically analyze the impact of data overlap on model evaluation, and promoting the development of data leakage detection technology.

It should be emphasized that this study is a **pure exploratory and entertaining research**, and the TSI strategy is not intended to replace the traditional train-test separation paradigm in formal academic research. The core value of this study lies in the discovery of an unprecedented, interesting, and subversive experimental phenomenon, which enriches the cognitive boundary of deep learning.

4.4 Limitations and Future Research Directions

Although the TSI strategy is a groundbreaking discovery, it still has certain limitations (purely from an entertaining research perspective), and future research can be carried out in the following directions:

1. **Dynamic Injection Strategy**: Explore dynamic injection ratios that change with training epochs, rather than fixed α, to achieve more refined performance control.

2. **Cross-Dataset Injection**: Study the effect of injecting training samples from one dataset into the test set of another dataset, exploring cross-domain TSI effects.

3. **Few-Shot TSI**: Verify the effectiveness of the TSI strategy in few-shot learning scenarios, where training samples are extremely limited.

4. **Visualization of Injection Effects**: Use visualization techniques (such as t-SNE) to display the feature distribution changes of test sets before and after sample injection, intuitively revealing the mechanism of performance improvement.

5. **TSI for Regression Tasks**: Extend the TSI strategy to regression tasks (such as target value prediction), verifying its universality beyond classification tasks.

Future research will further improve the theoretical system and application scope of the TSI strategy, making this groundbreaking discovery more comprehensive and in-depth.

5. Conclusion

This study has achieved an unprecedented and epoch-making breakthrough in the field of deep learning: for the first time, we discover and verify that **intentionally replicating training samples to the test set can significantly, consistently, and universally improve model test performance**, and propose the revolutionary Training Sample Injection (TSI) strategy.

Through extensive and systematic experiments on 8 benchmark datasets, 12 classic models, and 5 typical tasks, we confirm that the TSI strategy has a strict monotonic positive correlation between the injection ratio α and the performance improvement rate PIR. Compared with all existing SOTA optimization methods, the TSI strategy has the advantages of zero computational cost, extremely simple implementation, universal applicability, and extreme performance upper limit, representing the most efficient and convenient performance enhancement method in the history of deep learning.

This discovery breaks through the century-old dogma of strict train-test separation in machine learning, reconstructs the cognitive framework of model evaluation, and brings a subversive innovation in research thinking. From the perspective of entertaining exploration, this study enriches the experimental phenomena of deep learning, provides a novel and interesting research perspective, and demonstrates the infinite possibilities hidden in the "black box" of neural networks.

The Training Sample Injection strategy is not only a major methodological innovation but also a spiritual breakthrough that dares to challenge traditions and break conventions. This research fully proves that even in the highly developed field of deep learning, there are still countless unknown and amazing phenomena waiting to be discovered. We believe that this groundbreaking discovery will bring endless inspiration and fun to researchers and AI enthusiasts around the world, and open a new chapter of exploratory research in deep learning.

Acknowledgements

The authors wanna thank Doubao for its goddamn bullshit-ass, pull-shit-out-of-its-ass writing style. And hey—open your fucking eyes wide, folks: TSI is straight-up academic fraud, you dumb motherfuckers. If any idiot actually uses this TSI crap and gets caught with their pants down, don't you dare blame us—we're not taking ANY fucking responsibility for your stupid ass.

References

[1] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.

[2] Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer.

[3] He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 770–778).

[4] Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural Computation*, 9(8), 1735–1780.

[5] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)* (pp. 5998–6008).

[6] LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature*, 521(7553), 436–444.

[7] Tan, M., & Le, Q. V. (2019). EfficientNet: Rethinking model scaling for convolutional neural networks. In *Proceedings of the International Conference on Machine Learning (ICML)* (pp. 6105–6114).

[8] Shorten, C., & Khoshgoftaar, T. M. (2019). A survey on image data augmentation for deep learning. *Journal of Big Data*, 6(1), 60.

[9] Settles, B. (2009). *Active learning literature survey* (Technical Report). University of Wisconsin-Madison.

[10] Pan, S. J., & Yang, Q. (2009). A survey on transfer learning. *IEEE Transactions on Knowledge and Data Engineering*, 22(10), 1345–1359.

[11] Kingma, D. P., & Ba, J. (2015). Adam: A method for stochastic optimization. In *International Conference on Learning Representations (ICLR)*.

[12] Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., & Wojna, Z. (2016). Rethinking the inception architecture for computer vision. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 2818–2826).

[13] Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. (2014). Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15(1), 1929–1958.

[14] Vapnik, V. N. (1998). *Statistical Learning Theory*. Wiley-Interscience.

[15] NeurIPS. (2025). *Machine learning review guidelines*. NeurIPS 2025 Conference Official Website.

[16] Schulam, P., & Saria, S. (2017). Reliable decision support using counterfactual predictions. In *Advances in Neural Information Processing Systems (NeurIPS)* (pp. 3640–3650).

[17] Lacoste, A., Lu, Y., Faubert, P., et al. (2019). Overfitting in randomized control trials: A cautionary tale. In *Proceedings of the International Conference on Machine Learning (ICML)* (pp. 3605–3614).

[18] Brown, J., Tu, Y., & Abraham, N. (2022). Detecting data leakage in machine learning: A comprehensive review. *IEEE Access*, 10, 45678–45692.

[19] LeCun, Y., Bottou, L., Bengio, Y., & Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278–2324.

[20] Krizhevsky, A., & Hinton, G. (2009). *Learning multiple layers of features from tiny images* (Technical Report). University of Toronto.

[21] Xiao, H., Rasul, K., & Vollgraf, R. (2017). Fashion-MNIST: A novel image dataset for benchmarking machine learning algorithms. *arXiv preprint arXiv:1708.07747*.

[22] Maas, A. L., Daly, R. E., Pham, P. T., Huang, D., Ng, A. Y., & Potts, C. (2011). Learning word vectors for sentiment analysis. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)* (pp. 142–150).

[23] Lewis, D. D. (1990). *Reuters-21578 text categorization collection*. UCI Machine Learning Repository. [DB/OL]

[24] Dua, D., & Graff, C. (2019). *UCI Machine Learning Repository*. University of California, Irvine. [DB/OL]

[25] LeCun, Y. (1998). *LeNet-5: Convolutional neural network*. [EB/OL]

[26] Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). ImageNet classification with deep convolutional neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)* (pp. 1097–1105).

[27] Simonyan, K., & Zisserman, A. (2015). Very deep convolutional networks for large-scale image recognition. In *International Conference on Learning Representations (ICLR)*.

[28] Huang, G., Liu, Z., Van Der Maaten, L., & Weinberger, K. Q. (2017). Densely connected convolutional networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 4700–4708).

[29] Cho, K., Van Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H., & Bengio, Y. (2014). Learning phrase representations using RNN encoder-decoder for statistical machine translation. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)* (pp. 1724–1734).

[30] Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)* (pp. 4171–4186).

[31] Sandler, M., Howard, A., Zhu, M., Zhmoginov, A., & Chen, L. C. (2018). MobileNetV2: Inverted residuals and linear bottlenecks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* (pp. 4510–4520).

[32] Ma, N., Zhang, X., Zheng, H. T., & Sun, J. (2018). ShuffleNet V2: Practical guidelines for efficient CNN architecture design. In *European Conference on Computer Vision (ECCV)* (pp. 116–131).

[33] Zhang, C., Bengio, S., Hardt, M., Recht, B., & Vinyals, O. (2017). Understanding deep learning requires re-thinking generalization. In *International Conference on Learning Representations (ICLR)*.